

Operating Systems

2-7029110

Lecture 1 - Intro



Ariel University
Computer Science Department

Welcome to Operating Systems course!

- Who Am I ? (Where is Dr. Kogan ?)
- Contacts and reception hours
- Why this course?
- Duties

Welcome to Operating Systems course!

- Who am I ?

Arkady (אהרן) Gorodischer.

About 10 years of developing experience

- Contacts

Phone: 0546407773 (limited availability)

Email: arkady82@gmail.com

Reception hours: Sunday, 17:30 – 19:00

Welcome to Operating Systems course!

Why this course?

OS is a running environment and should be effectively used. It may help with a lot of already made tools but may become a nightmare if one is not aware of how should it be handled.

Welcome to Operating Systems course!

Duties:

- No attendance is required but what is told during the lectures is obligatory.
- 70% - Exam (> 56)
30% - Homework (avg > 56)

A little bit of History

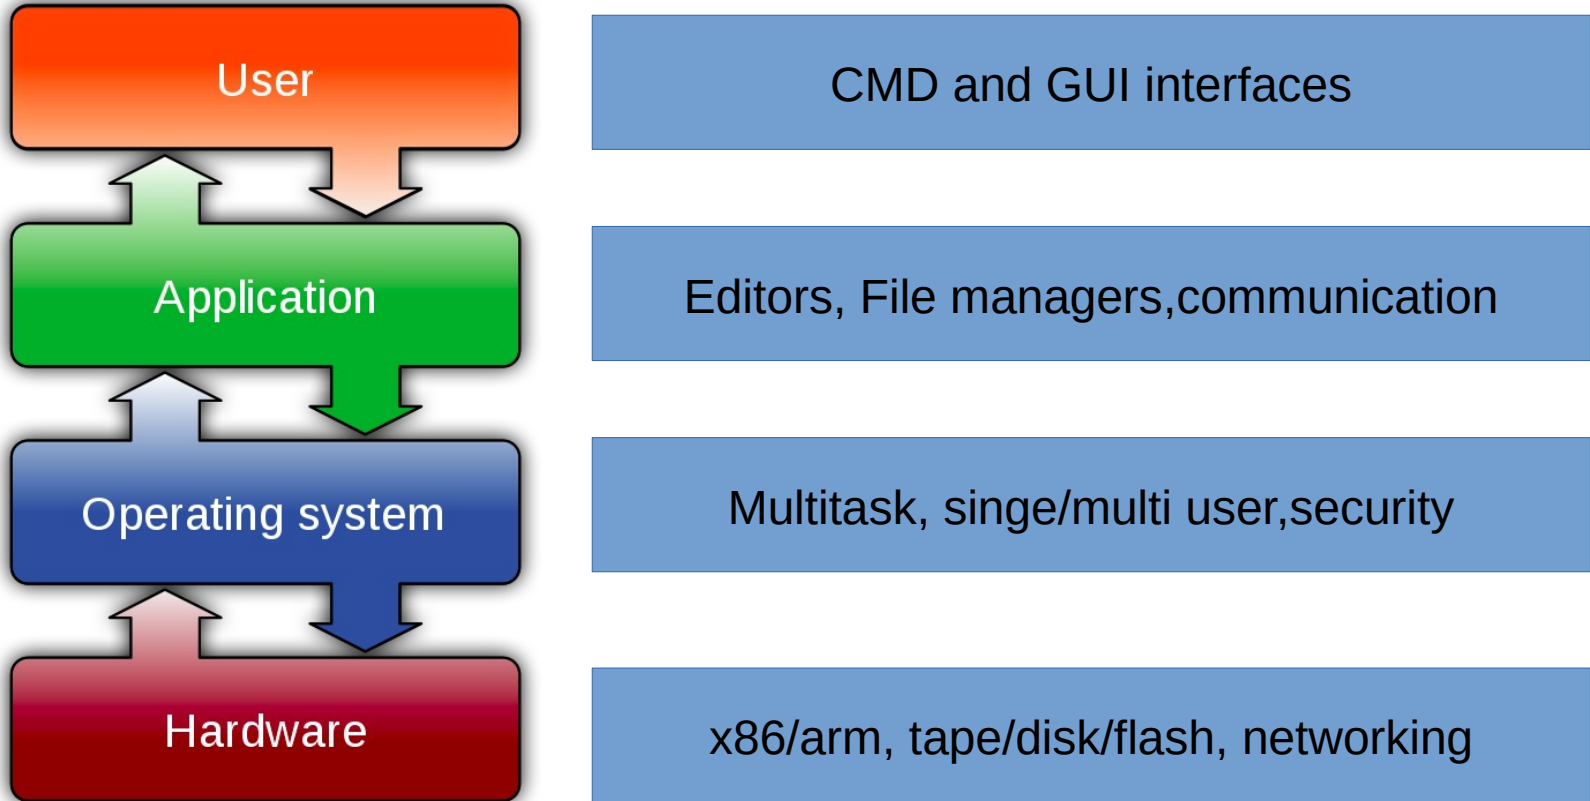
- 1960's – IBM OS/360 – OS for a product line
- 1970's – PLATO – Chat, multi-player games
- 1980's – MS DOS – for a PC. CMD , single user
- 1980's – Apple (Xerox) Windows interface
- 1990's – Win95,NT, GNU/Linux, Symbian
- 2000's - WinXP, iOS, Android
- 2010's – Raspbian, ROS, SailFish, Fuchisa
- 2020's – You may change the world :-)

What is OS - Operation System

The world before and after

- Bootloader /BIOS – minimalistic OS
- Bare Metal / Machine – how was it done
Extended Machine – the better way
- OS:
An Extended machine
A Resource Manager

Operation System Place



Extended Machine


- Provides
 - stable: doesn't crash
 - portable: can run code on more than one type of machine
 - reliable: always reacts in the same way
 - safe: doesn't do something dangerous
 - well-behaved: acts in a proper manner environment
- Computer “appears” to more than it is
 - “appears” to be many processors
 - “appears” to be many, large memories
- Features:
 - threads, processes, files, communication channels

Resource Manager

- Support many devices simultaneously
e.g. keyboard, mouse, printer, speakers, microphone
- Share resources among users and programmes
fairly: each programme gets a change to run
safely: protects against corruption
efficiently: using the available resources to provide the best service possible
- Allocates resources to users
Disks, memory, network interfaces, timers,
terminals/displays, laser printers, etc
Who's using what?
How is it shared?

Process management


- Process and Thread Management, Synchronization and Scheduling



CPU management


- Processing and I/O
- Computation communication overlapping
- Scheduler

Multi-Core Challenges

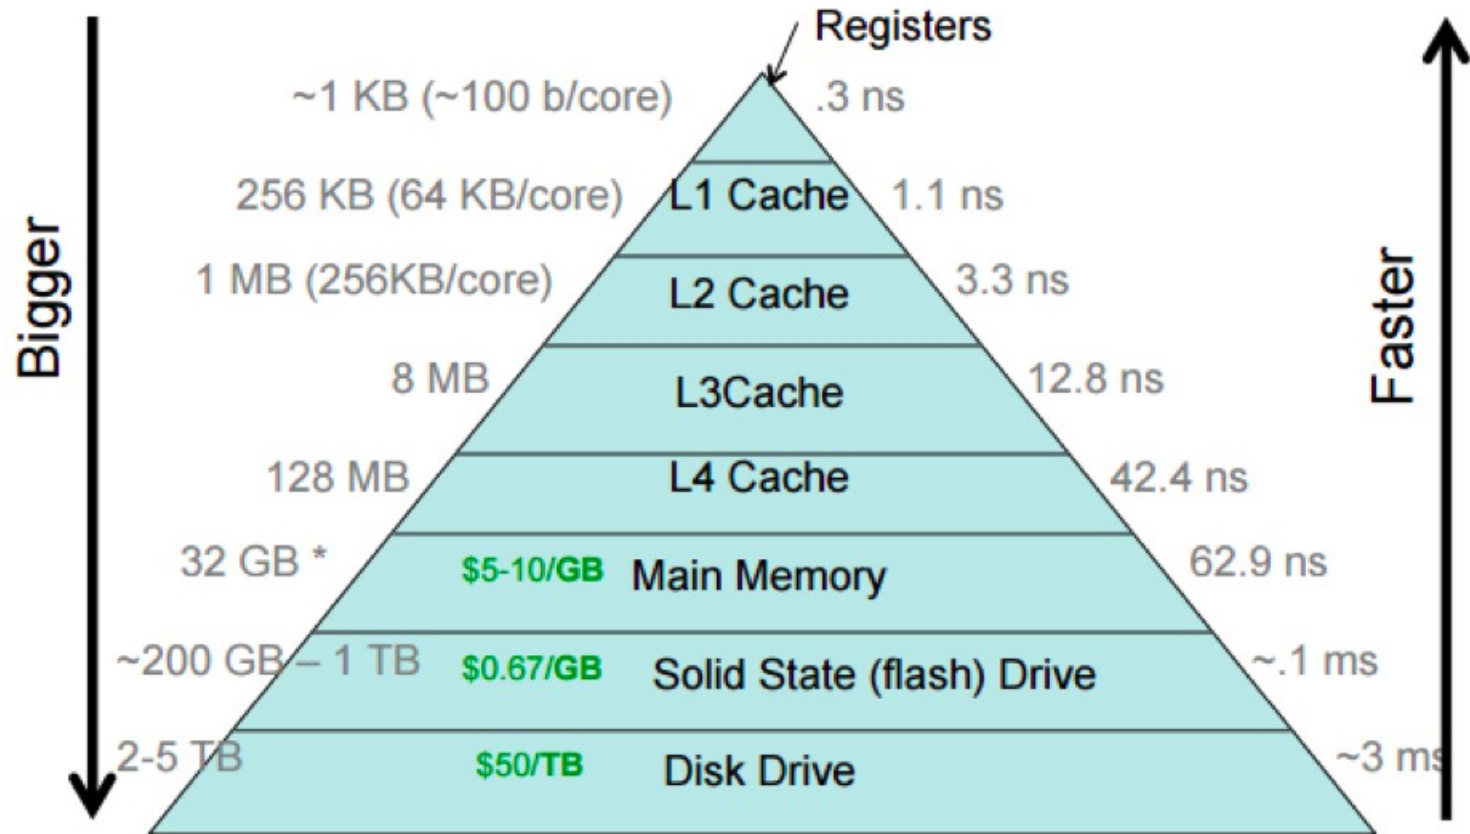


Memory management

- Main Memory management
זיכרון ראשי
- Secondary storage management
זכרון משני



Memory management



Interrupts and Interrupts handlers

```
Do forever{  
    IR = memory[PC];  
    execute(IR);  
    PC++;  
    If(Interrupt_Request) {  
        memory[0] = PC;  
        PC = memory[1]  
    }  
}
```

"User space" vs "Kernel space"

user space vs. kernel space

JULIA EVANS
@b0rk

drawings.jvns.ca

the Linux kernel has
millions of lines of code


- ★ read+write files
- ★ decide which programs get to use the CPU
- ★ make the keyboard work

When Linux kernel
code runs, that's
called

kernel space

When your program
runs, that's

user space



time for a
Context switch
to kernel space

your program switches
back and forth

```
str = "my string"
```

```
x = x + 2
```

```
file.write(str) ← ★ switch to kernel space ★
```

```
y = x + 4
```

```
str = str * y ← ★ and we're back to user space! ★
```

timing your process

```
$ time find /home
```

0.15 user 0.73 system

↑
time spent in
your process

↑
time spent by
the kernel doing
work for your
process

User “space” vs “Kernel space”

The division is made to protect the system, and separate sensitive system parts from user code

Protecting CPU Privileged instructions

Protecting OS Scheduler, Interrupts, etc

So, how can we use the kernel ??

System Calls

An application can interact with the kernel via system call (syscall). A Dedicated API, that cares more about security, than being a user-friendly.

Ex c++ fread() -> -> kernel read()

<https://linuxhint.com/what-is-a-linux-system-call/>

https://linuxhint.com/linux_system_call_tutorial_c/

System Calls 32bit example

```
# -----
# Writes "Hello, World" to the console using only system calls. Runs on 32-bit Linux only.
# To assemble and run:
#
#   gcc -m32 -c hello-32.s && ld hello-32.o && ./a.out
#
# or
#
#   gcc -m32 -nostdlib hello-32.s && ./a.out
# -----

.global _start

.text
_start:
    # write(1, message, 13)
    mov     $4, %eax           # system call 4 is write
    mov     $1, %ebx           # file handle 1 is stdout
    mov     $message, %ecx      # address of string to output
    mov     $13, %edx           # number of bytes
    int     $0x80              # invoke operating system to do the write

    # exit(0)
    mov     $1, %eax           # system call 1 is exit
    xor     %ebx, %ebx          # we want return code 0
    int     $0x80              # invoke operating system to exit
message:
    .ascii "Hello, world\n"
```

System Calls 64bit example

```
# -----  
# Writes "Hello, World" to the console using only system calls. Runs on 64-bit Linux only.  
# To assemble and run:  
#  
#   gcc -c hello.s && ld hello.o && ./a.out  
#  
# or  
#  
#   gcc -nostdlib hello.s && ./a.out  
# -----  
  
    .global _start  
  
    .text  
_start:  
    # write(1, message, 13)  
    mov    $1, %rax          # system call 1 is write  
    mov    $1, %rdi          # file handle 1 is stdout  
    mov    $message, %rsi    # address of string to output  
    mov    $13, %rdx         # number of bytes  
    syscall                  # invoke operating system to do the write  
  
    # exit(0)  
    mov    $60, %rax         # system call 60 is exit  
    xor    %rdi, %rdi        # we want return code 0  
    syscall                  # invoke operating system to exit  
message:  
    .ascii "Hello, world\n"
```